

Lecture 3

# Linear methods

Machine Learning  
Sergey Muravyov

29.01.2024

# Lecture plan

- Linear classification
- Gradient descent
- Linear regression and matrix decomposition
- Regularization

# Lecture plan

- Linear classification
- Gradient descent
- Linear regression and matrix decomposition
- Regularization

# Problem formulation

**Constraint:**  $Y = \{-1, +1\}$

$T^\ell = \{(x_i, y_i)\}_{i=1}^\ell$  is given

Find classifier  $a_w(x, T^\ell)$  in form  $\text{sign}(f(x, w))$ ,  
where  $f(x, w)$  is a discernment function,  
 $w$  is a parameter vector.

**Key hypothesis:** objects are (well-)separable.

**Main idea:** search among separating surfaces described  
with  $f(x, w) = 0$ .

# Margin

**Margin** of object  $x_i$ :

$$M_i(w) = y_i f(x_i, w),$$

$M_i(w) < 0$  is an evidence of misclassification.

# Margin

**Margin** of object  $x_i$ :

$$M_i(w) = y_i f(x_i, w),$$

$M_i(w) < 0$  is an evidence of misclassification.

We have previously defined **margin** of object  $x_i$  as

$$M(x_i) = C_{y_i}(x_i) - \max_{y \in Y \setminus \{y_i\}} C_y(x_i),$$

where  $C_y(u) = \sum_{i=1}^{\ell} [y(u, i) = y] w(i, u)$ ,  $w(i, u)$  is function of  $u$ 's  $i$ th neighbor importance.

What is their relation?

# Loss function smoothing

Empirical risk:

$$Q(a_w, T^\ell) = Q(w) = \sum_i^\ell [M_i(w) < 0],$$

is simply the number of errors classifier  $a_w$  shows.

The function is not smooth, so it is hard to find optima.


Approximation:

$$\tilde{Q}(w) = \sum_i^\ell L(M_i(w)),$$

where  $L(M_i(w)) = L(a_w(x_i, T^\ell), x_i)$  is a loss function.

# Smooth loss functions

We want  $L$  to be non-negative, non-increasing, and smooth:



$H(M) = (-M)_+$	— piecewise linear (Hebb's rule);
$V(M) = (1 - M)_+$	— piecewise linear (SVM);
$L(M) = \log_2(1 + e^{-M})$	— logarithmic (LR);
$Q(M) = (1 - M)^2$	— square (LDA);
$S(M) = 2(1 + e^M)^{-1}$	— sigmoid (ANN);
$E(M) = e^{-M}$	— exponential (AdaBoost).



# Linear classifier

$f_j: X \rightarrow \mathbb{R}, j = 1, \dots, n$  are numeric features.

**Linear classifier:**

$$a_w(x, T^\ell) = \text{sign} \left( \sum_{i=1}^n w_i f_i(x) - w_0 \right).$$

$w_1, \dots, w_n \in \mathbb{R}$  are feature **weights**.

Equivalent notation:

$$a_w(x, T^\ell) = \text{sign}(\langle w, x \rangle),$$


if a feature  $f_0(x) = -1$  is added.

# Neuron

**McCulloch-Pitts neuron:**

$$a_w(x, T^\ell) = \sigma \left( \sum_{i=1}^n w_i f_i(x) - w_0 \right),$$

where  $\sigma$  is an activation function.



# Set of algorithms

That assumption of how the classifier should look like specified the algorithm set  $A_{\text{linear}}$  from which we must choose an algorithm.

**How does this algorithm set look like?**

# Set of algorithms

That assumption of how the classifier should look like specified the algorithm set  $A_{\text{linear}}$  from which we must choose an algorithm.

How does this algorithm set look like?

$$A_{\text{linear}} = \{a_w(x) = \text{sign}(\langle w, x \rangle) | w \in \mathbb{R}^n\}$$

# How to learn a linear classifier?

We should find parameter vector  $w$ .

We can use almost any optimization algorithm capable to optimize empirical risk in the corresponding space.

The empirical risk is not black-box function.

Even more, we have guaranteed that it is a smooth function.

# Lecture plan

- Linear classification
- Gradient descent
- Linear regression and matrix decomposition
- Regularization

# Gradient descent

Empirical risk minimization problem

$$\tilde{Q}(w) = \sum_i^{\ell} L(M_i(w)) = \sum_i^{\ell} L(\langle w, x_i \rangle y_i) \rightarrow \min_w.$$

**Gradient descent (a.k.a Batch gradient descent):**

$w^{[0]}$  = **an initial guess value**;

$$w^{[k+1]} = w^{[k]} - \mu \nabla Q(w^{[k]}),$$

where  $\mu$  is **a gradient step** a.k.a **learning rate**.

$$w^{[k+1]} = w^{[k]} - \mu \sum_i^{\ell} L'(\langle w, x_i \rangle y_i) x_i y_i.$$

# Stochastic gradient descent

Problem is that there are too many objects function of which should be reestimated on each step.

**Stochastic gradient descent:**

$w^{[0]}$  is **an initial guess values**;

$x_{(1)}, \dots, x_{(\ell)}$  is **an objects order**;

$$w^{[k+1]} = w^{[k]} - \mu L'(\langle w^{[k]}, x_{(k)} \rangle y_{(k)}),$$

$$Q^{[k+1]} = (1 - \alpha)Q^{[k]} + \alpha L(\langle w^{[k]}, x_{(k)} \rangle y_{(k)}).$$

Stop when values of  $Q$  and/or  $w$  do not change much.



# Mini-batch gradient descent

Problem is that it is a bit too random because it depends only on a single object.

**Mini-batch gradient descent:**

$w^{[0]}$  is **an initial guess values**;  $b$  is **a batch size**

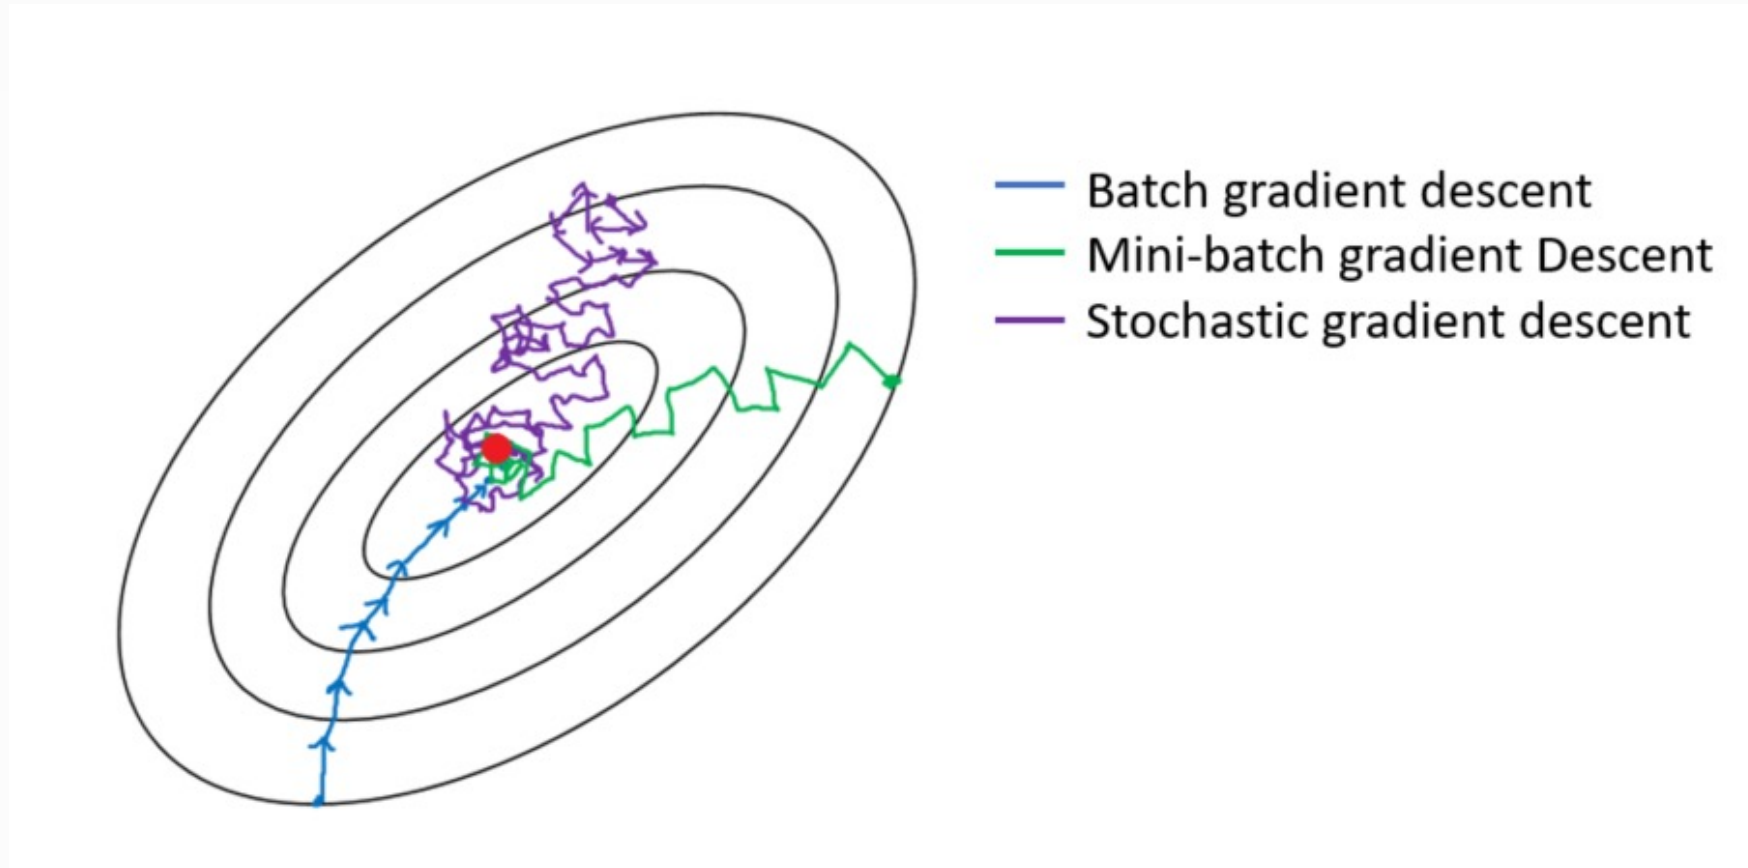
$x_{(1)}, \dots, x_{(\ell)}$  is **an objects order**;

$$w^{[K+1]} = w^{[K]} - \mu \sum_{k=Kb}^{k=(K+1)b} L'(\langle w^{[K]}, x_{(k)} \rangle y_{(k)}),$$

$$Q^{[K+1]} = (1 - \alpha)Q^{[K]} + \alpha \sum_{k=Kb}^{k=(K+1)b} L(\langle w^{[K]}, x_{(k)} \rangle y_{(k)}).$$

Stop when values of  $Q$  and/or  $w$  do not change much.

# Comparison



# Rosenblatt's rule and Hebb's rule

**Rosenblatt's rule** for  $\{1; 0\}$  classification case for weight learning: for each object  $x_{(k)}$  change the weight vector:

$$w^{[k+1]} := w^{[k]} - \eta(a_w(x_{(k)}) - y_{(k)}).$$

**Hebb's rule** for  $\{1; -1\}$  classification case for weight learning: for each object  $x_{(k)}$  change the weight vector:

if  $\langle w^{[k]} x_{(k)} \rangle y_{(k)} < 0$  then  $w^{[k+1]} := w^{[k]} + \eta x_{(k)} y_{(k)}$ .

# Novikov's theorem

## Theorem (Novikov)

Let sample  $T^\ell$  be linearly separable:  $\exists \tilde{w}, \exists \delta > 0$ :


$\langle \tilde{w}, x_i \rangle y_i > \delta$  for all  $i = 1, \dots, \ell$ .

Then the stochastic gradient descent with Hebb's rule will find weight vector  $w$ , which:

- splits sample without error;
- with any initial guess  $w^{[0]}$ ;
- with any learning rate  $\mu > 0$ ;
- independently on objects ordering  $x_{(i)}$ ;
- with finite numbers of changing vector  $w$ ;
- if  $w^{[0]} = 0$ , then the number of changes in vector  $w$  is

$$t_{\max} \leq \frac{1}{\delta^2} \max ||x_j||.$$

# Convergence problems




# Heuristics for initial guesses

Important for non-convex functions

- $w_j = 0$  for all  $j = 0, \dots, n$ ;
- small random values:  $w_j \in \left[-\frac{1}{2n}, \frac{1}{2n}\right]$ ;
- $w_j = \frac{\langle y, f_j \rangle}{\langle f_j, f_j \rangle}$ ;
- learn it with a small random subsample;
- multiple runs with different initial guesses.

# Learning rates comparison



# Heuristics for learning rate

- Convergence is achieved for convex functions when

$$\mu^{[k]} \rightarrow 0, \sum \mu^{[k]} = \infty, \sum (\mu^{[k]})^2 < \infty$$

- **Steepest gradient descent:**

$$Q(w^{[k]} - \mu^{[k]} \nabla Q(w^{[k]})) \rightarrow \min_{\mu^{[k]}}$$

- Steps for “jog of” local minima
- Second order methods
- Usage of mean vector of recent steps



# Heuristics for object ordering

- take objects from different classes each step;
- take misclassified objects more frequently;
- do not take “good” object, such that  $M_i > \kappa_+$ ;
- do not take noisy objects, such that  $M_i < \kappa_-$ .

# SG algorithm discussion

## Advantages:

- it is easy to implement;
- it is easy to generalize for any  $f$  and  $L$ ;
- dynamical learning;
- can handle small samples.

## Disadvantages:

- slow convergence or even divergence is possible;
- can stuck in local minima, saddle points;
- proper heuristic choice is very important;
- overfitting.

# Lecture plan

- Linear classification
- Gradient descent
- Linear regression and matrix decomposition
- Regularization

# Linear regression model

Model of multidimensional linear regression:

$$f(x, \theta) = \sum_{j=1}^n \theta_j f_j(x), \quad \theta \in \mathbb{R}^n.$$

Matrix notation:

$$F = \begin{pmatrix} f_1(x_1) & \dots & f_n(x_1) \\ \dots & \dots & \dots \\ f_1(x_\ell) & \dots & f_n(x_\ell) \end{pmatrix}, y = \begin{pmatrix} y_1 \\ \dots \\ y_\ell \end{pmatrix}, \theta = \begin{pmatrix} \theta_1 \\ \dots \\ \theta_n \end{pmatrix}.$$

Quality in matrix notation:

$$Q(\theta, T^\ell) = \sum_{i=1}^{\ell} (f(x_i, \theta) - y_i)^2 = \|F\theta - y\|^2 \rightarrow \min_{\theta \in \mathbb{R}^n}.$$

# Matrix decomposition

There are plenty of ways how it can be solved.

One of the most popular way is to apply singular **vector** decomposition, which is a **matrix decomposition** (a.k.a. **matrix factorization**) method.

# Normal equation system

Minimum condition:

$$\frac{\partial Q(\theta)}{\partial \theta} = 2F^T(F\theta - y) = 0.$$
$$\theta^* = (F^T F)^{-1} F^T y$$

$F^+ = (F^T F)^{-1} F^T$  is **pseudo reverse matrix (Moore–Penrose inverse)**

$P_F = FF^+$  is **projection matrix**

Solution:

$$\theta^* = F^+ y.$$

Minimum approximation:

$$Q(\theta^*) = ||P_F y - y||^2.$$

# Singular vector decomposition

**Theorem:** any matrix  $F$  size of  $\ell \times n$  can be represented with singular decomposition

$$F = VDU^{\top}.$$

With

- $V = (v_1, \dots, v_n)$  is size of  $\ell \times n$  and orthogonal  $V^{\top}V = I_n$ , columns  $v_j$  are eigenvectors of matrix  $FF^{\top}$ ;
- $U = (u_1, \dots, u_n)$  is size of  $n \times n$  and orthogonal  $U^{\top}U = I_n$ , rows  $u_j$  are eigenvectors of matrix  $F^{\top}F$ ;
- $D = \text{diag}(\sqrt{\lambda_1}, \dots, \sqrt{\lambda_n})$  is size of  $n \times n$ ,  $\sqrt{\lambda_j}$  are **singular numbers**, squares of eigenvalues of matrices  $FF^{\top}$  and  $F^{\top}F$ .

# SVD interpretation

You can think of some latent space, to which we want to project data.

$D$  represents importance of each basis vector

$V$  represents how objects correspond to the basis vectors

$U$  represents how features correspond to the basis vectors



# OLS with SVD

$$F^+ = (UDV^\top VDU^\top)^{-1}UDV^\top = UD^{-1}V^\top = \sum_{j=1}^n \frac{1}{\sqrt{\lambda_j}} u_j v_j^\top;$$

$$\theta^* = F^+ y = UD^{-1}V^\top y = \sum_{j=1}^n \frac{1}{\sqrt{\lambda_j}} u_j (v_j^\top y);$$

$$F\theta^* = P_F y = (VDU^\top)UD^{-1}V^\top y = VV^\top y = \sum_{j=1}^n v_j (v_j^\top y);$$

$$\|\theta^*\|^2 = \|D^{-1}V^\top y\|^2 = \sum_{j=1}^n \frac{1}{\lambda_j} (v_j^\top y)^2.$$

# Discussion


- When we can compute SVD, we can easily find solution for OLS.
- SVD computations are quite heavy still, its complexity is  $O(\ell^2 n + n^3)$
- SVD is important in many other machine learning tasks, first of all, in dimensionality reduction.

# Lecture plan

- Linear classification
- Gradient descent
- Linear regression and matrix decomposition
- Regularization

# Overfitted algorithm (reminder)

$$y(x) = \frac{1}{1 + 25x^2}; \quad a(x) \text{ — polynomial of degree } n = 38$$



# Regularization

**Key hypothesis:**  $w$  “swings” during training causing overfitting

**Main idea:** clip  $w$  norm.

Add regularization penalty for weights norm:

$$Q_{\tau}(a_w, T^{\ell}) = Q(a_w, T^{\ell}) + \frac{\tau}{2} \|w\|^2 \rightarrow \min_w.$$

$\tau$  is a coefficient representing the strength of model regularization, or, equally, is a trade-off between model performance and generalizability.

# Ridge regression

**Assumption:** values of  $\theta$  have Gaussian distribution with covariance matrix  $\sigma I_n$ :

$$Q_\tau(\theta) = ||F\theta - y||^2 + \frac{1}{2\sigma} ||\theta||^2 \rightarrow \min_{\theta},$$

where  $\tau = 1/\sigma$  is regularization penalty.

It can be simply added to OLS solution:

$$\theta_\tau^* = (F^\top F + \tau I_n)^{-1} F^\top y.$$

# Solution for ridge regression

$$\theta_{\tau}^* = U(D^2 + \tau I_n)^{-1} D V^{\top} y = \sum_{j=1}^n \frac{\sqrt{\lambda_j}}{\lambda_j + \tau} u_j (v_j^{\top} y);$$

$$F \theta_{\tau}^* = (V D U^{\top}) \theta_{\tau}^* = V \operatorname{diag} \left( \frac{\lambda_j}{\lambda_j + \tau} \right) V^{\top} y =$$

$$= \sum_{j=1}^n \frac{\lambda_j}{\lambda_j + \tau} v_j (v_j^{\top} y);$$

$$\|\theta^*\|^2 = \|D^2 (D^2 + \tau I_n)^{-1} D^{-1} V^{\top} y\|^2 = \sum_{j=1}^n \frac{1}{\lambda_j + \tau} (v_j^{\top} y)^2.$$

# Tibshirani lasso

**Assumption:** values of vector  $\theta$  has Laplacian distribution:

$$\begin{cases} Q_{\tau}(\theta) = ||F\theta - y||^2 \rightarrow \min_{\theta}; \\ \sum_{i=1}^n |\theta_i| \leq \kappa. \end{cases}$$

**LASSO** (least absolute shrinkage and selection operator). Will lead to feature selection.



# LASSO regression

The resulting optimization problem

$$Q_{\tau}(\theta) = ||F\theta - y||^2 + \tau||\theta||_1 \rightarrow \min_{\theta},$$

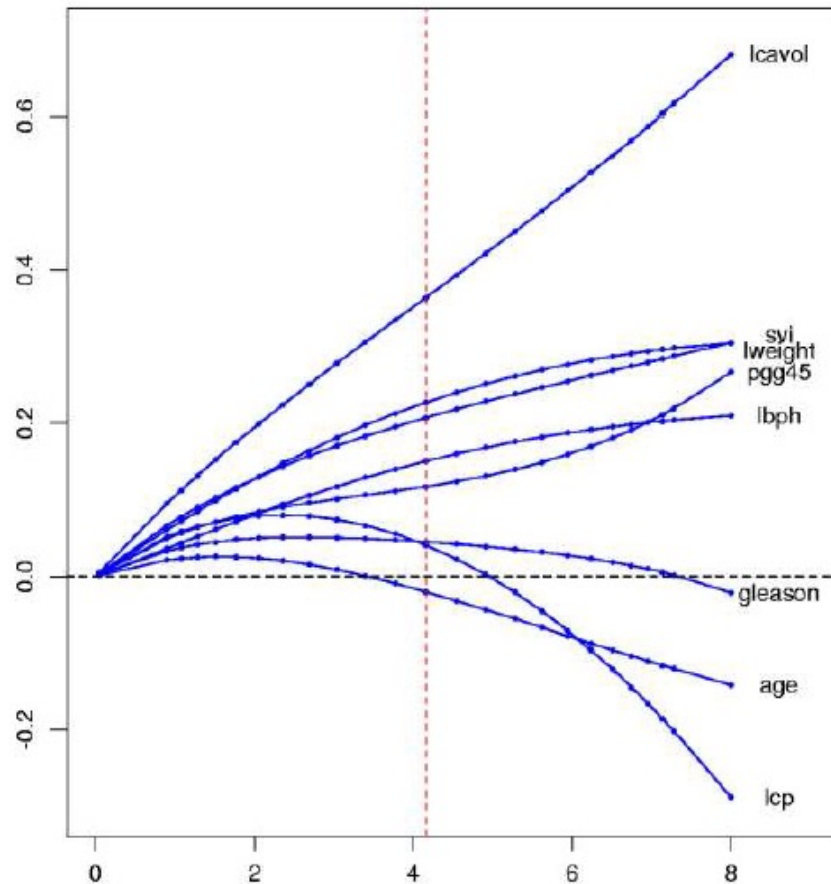
where  $||\theta||_1$  is  $l_1$ -norm:  $||\theta||_1 = \sum |\theta_i|$ .

No nice analytical solution exist.


However, a nice computational solution exist.

# Comparison

## Ridge regression



## Lasso



# Regularization for gradient descent

Regularization is simply adopted to any gradient descent:

$$\begin{aligned}\nabla Q_\tau(w) &= \nabla Q(w) + \tau w, \\ w^{[k+1]} &= w^{[k]}(1 - \mu\tau) - \mu\nabla Q(w).\end{aligned}$$

# Regularization discussion (1/2)

- $l_1$ -norm and  $l_2$ -norm regularizers are the most popular
- **ElasticNet**, which is sum of the previous two is also popular
- Many other may be used with respect to initial assumptions
- Some techniques are de-facto regularization or can be interpreted as regularization

# Regularization discussion (2/2)

- One of two general approaches to fight overfitting
- Required to solve ill-stated problems
- Regularization penalties can be simply added to represent properties you expect from model
- Regularization coefficient should be tuned
- Not always clean, how to choose regularization